

UPDATE REPORT · 18 SEPTEMBER 2026 · 21:10 PT

JARVIS gap report: job page extraction vs. your spec

This is the exact distance between what you asked for in the 14-page **Job Page Extraction Fix** and what is in the running build. It also carries forward unfinished proof from the earlier Comparison Result Fix. Nothing here is marked done without microphone proof.

Build: restarted after this update | **Tests:** 178 automated passed | **Live Deloitte proof:** not yet passed

1. The gap in one paragraph

The comparison workflow now runs and no longer answers “Completed: Compare and list gaps.” That was the previous bug. The new live failure was different: JARVIS compared your resume to **Chrome chrome** — Zoom 110%, Bookmark this tab, Install Reddit, Energy Saver, Managed bookmarks, Ask Google — because page text came from a flat accessibility dump of the whole Chrome HWND. This update stops that dump. Job analysis now prefers the actual document at the active tab URL, JSON-LD JobPosting first, then qualification sections, and it rejects toolbar / extension labels. That is shipped in code and covered by unit tests. It is **not** live-certified on the Deloitte Agentic AI Engineer, Senior (359035) posting. Until that exact live test prints WEB_DOCUMENT provenance and zero browser-UI requirements, the extraction spec is still open.

2. Scoreboard — you vs. this build

Your requirement	What the build does now	Status
Hard boundary: BROWSER_CHROME vs WEB_DOCUMENT	PageModel scopes exist. Job extraction accepts WEB_DOCUMENT only. Full-window UIA dump is no longer the job-text source.	SHIPPED
Source provenance on every compared requirement (scope, method, section)	Each RequirementItem stores source_scope, source_method, job_source_section, job_source_text. Logged per item.	SHIPPED
DOM-first source order: extension bridge → CDP → document UIA → page reader → vision	Page-reader HTTP fetch of the active URL + JSON-LD is primary. Document-subtree UIA is fallback. No Chrome extension. No CDP attach. No vision fallback.	PARTIAL
Do not use the top-level Chrome accessibility tree as page text	get_page_text() no longer calls collect_window_text(). Document UIA walks only the Document control.	SHIPPED
JSON-LD @type=JobPosting as highest-confidence source	Implemented. Confirmed against the live Deloitte HTML in a probe: JobPosting present, qualifications + description parseable.	SHIPPED
JobPostingModel with title, company, job_id, sections, skills, travel, sponsorship, clearance, tools	Model exists and fills title/company/job_id/required/preferred/education/travel/sponsorship from JSON-LD + section parser. Salary/location/employmentType only if JSON-LD has them.	PARTIAL
Section-aware extraction; responsibilities are context, not automatic requirements	Required / Preferred / Qualifications drive comparison. Responsibility sections are classified RESPONSIBILITY and not auto-added.	SHIPPED

Your requirement	What the build does now	Status
JobRequirementValidator rejects toolbar, extensions, nav, URLs, permission prompts	Reject list includes every forbidden example from the spec (Zoom, Bookmark, Reddit, Energy Saver, Ask Google, side panel, site information, etc.).	SHIPPED
Abort comparison if any requirement is BROWSER_CHROME / EXTENSION / OS_UI	Returns JOB_CONTENT_CONTAMINATION_DETECTED and no comparison list. User is told the listing was mixed with browser controls.	SHIPPED
Content contamination score; rebuild if over threshold	Score is computed. High score without JSON-LD drops page_type confidence. It does not yet rebuild from a second DOM path.	PARTIAL
page_type confirmed from evidence, not URL/title alone	JOB_LISTING requires structured data / qualifications / job id / apply evidence. JOB_APPLICATION looks for form/upload/required fields. Not live-proven.	PARTIAL
User-facing output contains real qualifications, never browser UI	Formatter skips noise labels. Unit tests assert Zoom/Bookmark/Reddit/Energy Saver never appear. Live Deloitte output not yet captured.	UNPROVEN
“What are the gaps?” from the clean stored ComparisonResult, planner_calls=0	RecentAnalysisFollowupResolver still in place from the previous update. Follow-ups will only be clean if the stored result is clean.	PARTIAL
Live Deloitte 359035: WEB_DOCUMENT for every compared requirement	Required. Not run on the restarted build yet.	GAP
Chrome extension / content-script DOM bridge	Not built. Spec listed this as preferred source #1.	GAP
DevTools / CDP attachment to existing Chrome	Not built. Spec listed this as preferred source #2.	GAP
Targeted vision fallback for document content	Not built. Spec listed this as source #5.	GAP

3. What I changed in this update

Root cause of the live failure: **collect_window_text(hwnd)** walked the entire Chrome window tree (toolbar, tab strip, extensions, side panel, plus the page) and handed 2,712 characters to JobResumeComparisonEngine. The engine then treated almost every line as a requirement. That produced 38 items, including Chrome UI.

Correct path now:

- Active tab URL from existing Chrome (already resolved; no Playwright about:blank).
- HTTP fetch of that URL (page-reader). Deloitte 359035 returns JSON-LD JobPosting.
- Parse JobPosting title, hiringOrganization, identifier, qualifications, description.
- HTML-to-text, then section split on Qualifications / Required / Preferred.
- If fetch is thin: UIA **Document** subtree only — never the top-level Chrome tree.
- JobRequirementValidator. Only VALID_REQUIREMENT and PREFERRED_REQUIREMENT reach compare.
- If chrome/extension labels remain and there is no clean JobPosting: abort, do not compare.

Files added or materially changed:

- `src/agent/browser/page_model.py` — scopes, page types, contamination score.
- `src/agent/browser/document.py` — fetch, JSON-LD, HTML text, document-only UIA.
- `src/agent/phase6/job_posting.py` — JobPostingModel + JobRequirementValidator.
- `src/agent/browser/existing.py` — `get_page_text` uses the document model.
- `src/agent/screen/uia.py` — `collect_document_text()` for renderer subtree only.
- `src/agent/phase6/compare.py` — compare from JobPostingModel; abort if contaminated.
- `src/agent/router.py` — “What does this page say?” → `browser.page_about`.

4. Carried-forward gaps from earlier specs

These were already implemented in code during the Comparison Result Fix. They are still open as **live proof**, and the chrome-contamination failure means the first live comparison cannot be counted as a pass.

Earlier requirement	Code	Live proof
Final answer is the actual comparison, never “Completed: Compare and list gaps.”	SHIPPED	UNPROVEN — last live run was contaminated, not a valid pass
HUD detailed list + short TTS summary	SHIPPED	UNPROVEN
“Okay, what are the gaps?” <code>planner_calls=0</code> from stored result	SHIPPED	UNPROVEN — will only be valid after a clean first compare
“What matched well?” / “Why did you say that one was covered?” from stored evidence	SHIPPED	UNPROVEN
Switch listing → “Compare this one to my resume” is a new Phase 6 task	SHIPPED	UNPROVEN
Wording: “not shown on your resume,” never “you do not have this”	SHIPPED	UNPROVEN on a clean Deloitte run

5. Required live tests — still yours to run

I cannot mark this update fixed. The spec is explicit: do not declare done until the exact Deloitte live test passes with `WEB_DOCUMENT` provenance on every compared requirement. JARVIS is restarted and waiting for these.

Test	What you do	Pass looks like	Current
A — Deloitte job	Open Agentic AI Engineer, Senior / 359035. Say the full compare command.	Real quals: 4+ AI/ML, Python, LangChain/LangGraph/AutoGen, degree, travel 50%. Zero Zoom/Bookmark/Reddit/Energy Saver.	GAP
B — Provenance	After A, check logs.	Every item: <code>source_scope=WEB_DOCUMENT</code> . <code>structured_jobposting_found=true</code> if JSON-LD present.	GAP
C — JSON-LD	Confirm the log line after A.	<code>structured_jobposting_found=true</code> and <code>structured_fields_used</code> includes <code>qualifications/description/title</code> .	GAP

Test	What you do	Pass looks like	Current
D — Extensions on	Leave Reddit/VPN/Energy Saver enabled. Repeat A.	Extension labels do not enter JobPostingModel or the spoken/HUD comparison.	GAP
E — Generic page	Open a non-job page. Ask “What does this page say?”	Summarizes the web document only. No toolbar/extensions.	GAP
F — Application form	Move from listing to the apply form.	page_type=JOB_APPLICATION with form fields / required / upload — not a fake requirements list.	GAP

6. What a passing Deloitte comparison should contain

Pulled from the live posting's JSON-LD during this update's probe — this is the job text JARVIS should be matching against, not Chrome UI:

Required (examples from the actual listing)

- 4+ years delivering AI/ML, with ≥ 1 year on Generative AI, Agentic AI, or multi-agent systems
- 2+ years hands-on building AI/ML solutions using Python
- 1+ year developing agentic systems (orchestration, tool use, autonomous workflows) — LangChain, Semantic Kernel, AutoGen, Strands, CrewAI, LangGraph
- Cloud agent stack across Azure and/or AWS and/or GCP
- 1+ year leading workstreams / translating business problems into AI solutions
- Bachelor's or Master's in a related technical field
- Ability to travel up to 50%
- Limited immigration sponsorship may be available

Preferred (examples)

- Consulting / client-facing delivery
- Multi-agent systems delivery
- Cloud/AI certifications
- LLM prompt engineering, fine-tuning, RAG architecture
- MLOps / AIOps, CI/CD, cloud-native development
- API design, microservices, event-driven systems

Travel, sponsorship, and certification status should land in **Needs confirmation** unless the resume actually states them. Responsibilities under “Work you’ll do” must not be dumped in as missing skills.

7. Known remaining engineering gaps (honest)

- **No in-page DOM bridge.** Existing Chrome is read by fetching the same URL JARVIS already resolved from the address bar. That is the real document for Deloitte (JSON-LD is in the HTML). It will miss pages that are login-walled, fully client-rendered with no JSON-LD, or different from what a logged-in session sees.
- **No CDP.** If a future listing is a SPA without JobPosting JSON-LD, fetch may return a shell. Document UIA is the only remaining path.
- **No vision fallback.** If both fetch and document UIA are thin, comparison should fail closed — it must not fall back to toolbar text.
- **Contamination score does not rebuild.** Spec says reject and rebuild from DOM/document. We reject/abort. We do not yet retry a second extractor.

- **JOB_APPLICATION is heuristic.** Form-field detection is from document text, not a real form control inventory. Test F may still miss upload/required flags.
- **Generic “what does this page say?”** now uses document text, but I have not watched it on a chrome-heavy non-job page.
- **Skill-name fallback** still scans clean document text for known skills (python, langchain, ...). That is only after chrome filtering. It must never scan a raw HWND dump.

8. Definition of done — still not met

Done for this spec means all of the following are true on your machine, on the real Deloitte tab, with extensions still enabled:

- `page_type=JOB_LISTING` from listing evidence, not from the URL string alone.
- `JobPostingModel` built; `structured_jobposting_found=true` on this posting.
- browser chrome excluded; extension nodes excluded.
- Every compared requirement `source_scope=WEB_DOCUMENT`.
- HUD/TTS contain real qualifications and zero forbidden terms (Zoom, Bookmark, Reddit, Energy Saver, Ask Google, side panel, site information).
- “What are the gaps?” returns only missing / partial / needs-confirmation from that clean model, `planner_calls=0`.

Until Test A and Test B pass, treat extraction as **open**. The old comparison-result delivery spec is also still open for live proof, because the only live comparison so far was invalid input.

9. How I will report after every later update

You asked for an in-depth PDF after every update so the remaining gap is never fuzzy. Going forward, each JARVIS change set will get a dated file under **docs/reports/** with the same three columns: your requirement, what the build does, and SHIPPED / PARTIAL / UNPROVEN / GAP. Live tests stay GAP until you run them. I will not collapse “code written” into “done.”

This report’s file: `docs/reports/JARVIS_GAP_REPORT_2026-09-18_job_page_extraction.pdf`

Generated 2026-09-18 21:12 PT from the job-page extraction update. Automated tests: 178 passed. Live certification: not claimed.